

Lesson 07.4: Generalized Linear Models

Instructor: Dr. GP Saggese, gsaggese@umd.edu

References:

- AIMA (Artificial Intelligence: a Modern Approach)
 - Chap 15: Probabilistic programming
- Martin, Bayesian Analysis with Python, 2018 (2e)




- ***Generalized Linear Models***
 - Simple Linear Model
 - Logistic Regression
 - Multiple Linear Regression

- Generalized Linear Models
 - *Simple Linear Model*
 - Logistic Regression
 - Multiple Linear Regression

Linear Model

- Many problems can be formulated as **linear regression**:
 - X and Y : uni-dimensional continuous RVs
 - Dataset of paired observations: $\{(x_1, y_1), \dots, (x_n, y_n)\}$
 - X : independent variables
 - Model/predict dependent variable Y
 - Assume linear relationship between Y and X

- E.g.,
 - Nobel laureates in a country vs amount of chocolate consumed
 - Sugar intake vs Blood glucose
 - Years of education vs Annual income
 - Advertising spend vs Sales
 - Study hours vs Exam score
 - Coffee consumption vs Productivity



Linear Model: Frequentist Approach

- A **linear model** is described by:

$$y = \alpha + \beta x$$

- **Frequentist approach:**

- Find parameters α, β using least square fitting
 - * Minimize average quadratic error between observed y and predicted $\hat{y}$

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (\alpha + \beta x_i - \hat{y}_i)^2$$

- Point-estimate from least squares corresponds to maximum a-posteriori with flat (uniform) priors on α and β
- With certain assumptions, it's “optimal”
 - * Residuals are Gaussian
 - * Linearity in parameters
 - * Independence
 - * Homoscedasticity (constant variance of ε_i)

Linear Model: Bayesian Approach

- **Bayesian assumptions:**

- Data y is Gaussian with mean $\alpha + \beta x$ and standard deviation

$$\sigma : y \sim N(\mu = \alpha + \beta x, \sigma)$$

- α, β, σ are independent
- Priors:

$$\begin{cases} \alpha \sim \text{Normal}(\mu_\alpha, \sigma_\alpha) \\ \beta \sim \text{Normal}(\mu_\beta, \sigma_\beta) \\ \sigma \sim \text{HalfNormal}(0, \sigma_\epsilon) \end{cases}$$

- Use **vague priors** to add problem knowledge:
 - Center intercept α around 0
 - Info on sign of β sometimes available
 - Use Half-Cauchy / exponential distribution for σ
 - Use uniform prior if parameter has hard boundaries

Linear Model: Synthetic Example

- Generate **random data from ground truth**

- $\alpha = 2.5$
- $\beta = 0.9$
- Add noise

```
[122]: np.random.seed(1)


# Number of samples.
N = 100

# Parameters.
alpha_real = 2.5
beta_real = 0.9
sigma_eps_real = 0.5

# Generate data.
x = np.random.normal(10, 1, N)
y_real = alpha_real + beta_real * x

# Add noise.
eps_real = np.random.normal(0, sigma_eps_real, size=N)
y = y_real + eps_real

[123]: plt.scatter(x, y_real)
plt.scatter(x, y);
```



Linear Model: Synthetic Example

- Create **linear model in PyMC**
- Use **vague priors**
 - Prior of intercept α and β centered around 0
 - Half-Cauchy distribution for σ

```
[125]: with pm.Model() as model_g:
        alpha = pm.Normal("alpha", mu=0, sigma=10)
        beta = pm.Normal("beta", mu=0, sigma=1)
        sigma = pm.HalfCauchy("sigma", 5)
        #
        mu = pm.Deterministic("mu", alpha + beta * x)
        y_pred = pm.Normal("y_pred", mu=mu, sigma=sigma, observed=y)
        idata_g = pm.sample(2000, tune=1000)
```


Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [alpha, beta, sigma]

Sampling 4 chains, 0 divergences

Sampling 4 chains for 1_000 tune and 2_000 draw iterations (4_000

```
[127]: pm.model_to_graphviz(model_g)
```


```
[127]:
```



Linear Model: Synthetic Example

- Run the sampler
- Ground truth
 - $\alpha = 2.5$
 - $\beta = 0.9$
- Recovered values
 - $\alpha = 2.12$
 - $\beta = 0.95$

```
[129]: az.plot_trace(idata_g, var_names=["alpha", "beta", "sigma"]);
```




```
[128]: az.summary(idata_g, var_names="alpha beta sigma".split(), kind="stats")
```

```
[128]:
```

	mean	sd	hdi_3%	hdi_97%
alpha	2.118	0.538	1.099	3.136
beta	0.946	0.053	0.845	1.046
sigma	0.476	0.034	0.412	0.538

Linear Model: Bike Rental Example

- Model relationship between:
 - **Temperature** X
 - **Number of bikes rented** Y



- Intermediate variable is mean number of bikes rented for temperature X :

$$\mu = \alpha + \beta X$$


$$Y \sim N(\mu, \sigma)$$

Linear Model: Bike Rental Example

- **Fit model** with PyMC:

$$\mu = 69 + 7.9X$$

- **Interpreting the model**
 - Temperature 0: expected rented bikes = 69
 - Each degree increase: expected rented bikes = +7.9
- **Parameters have uncertainty**
 - Posterior accounts for combined uncertainty
 - Bands represent quantiles [0.25, 0.75] and [0.03, 0.97] of prediction



- **Do you see any problem in the model?** 

Linear Model: Bike Rental Example (Criticism)

- **Problems**
 1. Model outputs negative bike numbers
 2. Model predicts real numbers, but count is discrete
- **Root cause:** Using Gaussian likelihood:
 - Extends to negative numbers
 - Is continuous
- **Solutions:**
 - *Hack*: Clip predictions below 0 and discretize output
 - *Elegant*: Use model defined for discrete positive numbers

Count Data

- **Count data** when random variable is discrete, bounded at 0
 - E.g., number of rented bikes
 - Large numbers use continuous distribution
 - * E.g., Gaussian
 - Sometimes modeled as discrete
 - * E.g., Poisson, negative binomial

Generalized Liner Model (GLM)

- **Generalization of linear model** using different distributions for likelihood:

$$\left\{ \begin{array}{l} \alpha \sim \text{prior} \\ \beta \sim \text{prior} \\ \mu = \alpha + \beta X \\ \theta \sim \text{prior} \\ Y \sim \phi(f(\mu), \theta) \end{array} \right.$$

where:

- ϕ : arbitrary distribution
 - * E.g., Normal, Student's t, negative binomial
- θ : auxiliary parameter for ϕ
 - * E.g., σ for Normal
- $f()$: “inverse link function” transforming μ from real line to domain of ϕ

Poisson Distribution

■ Poisson distribution


- Models events in a fixed interval (time, space)
- Assumes independent events at a constant rate

■ Examples

- *Call centers*: customers per hour
- *Natural events*: earthquakes in a region
- *Traffic flow*: cars through a toll booth per day
- *Biology*: mutations in a DNA segment

■ Cons

- Assumes mean equals variance (can't model "overdispersion")




Negative Binomial Distribution

- **Negative binomial distribution**
 - Models failures/trials for fixed successes in IID Bernoulli trials
 - Generalizes geometric distribution, modeling trials to first success
 - Models overdispersion

mean $\ll$ variance

- **E.g.,**
 - *Customer service*: unsuccessful interactions before success
 - *Sports*: games lost before winning fixed number of games





Overdispersion

- **Overdispersion** occurs when variance of count data is much larger than mean (e.g., 15-20x)
- **Examples:**
 - *Epidemiology*: modeling disease outbreak, number of new infections fluctuates widely
 - *Customer service*: large variation in daily complaints
- **How to model it?**
 - *Use Poisson*:
 - * **✗**: mean and variance must be equal
 - *Use Gaussian*:
 - * Mismatch predicting negative rented bikes
 - * Poor fit on positive side
 - *Use Negative Binomial*:
 - * Better fit, though not perfect
 - * Right tail predictions differ, but high demand probability is low
 - * Second parameter controls variance
 - * **✓**: overall better than Normal model

Robust Regression

- **Outliers** pull regression line away from data
- **Robust regression** is a generalized linear model
 - Avoids feature transform/winsorization
- Student's t-distribution
 - Heavier tails than Normal
 - Reduces outlier impact



- Generalized Linear Models
 - Simple Linear Model
 - *Logistic Regression*
 - Multiple Linear Regression

Logistic Regression


- **Logistic regression model**
 - Generalized linear model
 - Models binary response variable
 - * E.g., ham/spam, safe/unsafe, cloudy/sunny, hotdog/not hotdog
- **Logistic model** is:

$$\begin{cases} \theta = \text{logit}(\alpha + \beta x) \\ y \sim \text{Bernoulli}(\theta) \end{cases}$$

- **Logistic function** (aka sigmoid)

$$\text{logit}(z) = \frac{1}{1 + \exp^{-z}}$$

- Convert real numbers from θ to $[0, 1]$ for Bernoulli distribution



Iris Dataset

- Classical dataset of flower measurements from 3 Iris species
 - `setosa`, `virginica`, `versicolor`

	<code>sepal_length</code>	<code>sepal_width</code>	<code>petal_length</code>	<code>petal_width</code>	<code>species</code>
0	5.1	3.5	1.4	0.2	<code>setosa</code>
1	4.9	3.0	1.4	0.2	<code>setosa</code>
2	4.7	3.2	1.3	0.2	<code>setosa</code>
3	4.6	3.1	1.5	0.2	<code>setosa</code>
4	5.0	3.6	1.4	0.2	<code>setosa</code>

- E.g., predict probability of a flower being `setosa` from `sepal_length`

Classification with Logistic Regression

- **Logistic regression:**

- Computes probability of output being a certain value
- Models $\theta = \Pr(Y = 1|X)$
- Classifies output using decision rule, e.g.,

$$\Pr(Y = \text{versicolor} \mid \text{sepal_length}) > 0.5$$

- Same treatment as frequentist approach, has Bayesian flavor
 - Once Bayesian, generalize more
 - No gradient descent, solved with Bayesian inference
- **Classification with logistic regression might seem a misnomer**
 - **Regression** predicts continuous variable
 - **Classification** predicts discrete variable
 - Is “regression” since it computes probability of output being a certain value

Boundary Decision for a Classifier

- **Boundary decision** δ for a classifier are the values of independent variables making output probability 0.5
- E.g., for logistic regression δ :

$$0.5 = \text{logit}(\alpha + \beta\delta)$$

$$0 = \alpha + \beta\delta$$

$$\delta = -\frac{\alpha}{\beta}$$

- Decision boundary uncertainty due to α and β uncertainty
- Probability threshold 0.5 for equal **misclassification risk**
 - Misclassification cost may not be symmetrical
 - Minimize false negatives
 - * E.g., “*patient has disease, disease missed*”
 - Minimize false positives
 - * E.g., “*patient doesn't have disease, predicted*”

Odds

- The **odds of event** “ $y = 1$ ” is the ratio of favorable to unfavorable events

$$\text{odds} \triangleq \frac{\Pr(y = 1)}{1 - \Pr(y = 1)}$$

- Transformation between probability, odds, and log-odds
 - More intuitive in “gambling”
 - E.g., odds of rolling a 2 on a fair die are $\frac{1/6}{5/6} = \frac{1}{5} = 0.2$
 - * One favorable event for 5 unfavorable events
- Interpreting **logistic regression in terms of odds**:
 - Since $\theta = \text{logit}(\alpha + \beta x)$ and $\theta = \Pr(y = 1)$:

$$\alpha + \beta x = \log\left(\frac{\Pr(y = 1)}{1 - \Pr(y = 1)}\right)$$




- Logistic regression models log-odds linearly
- β is the increase in log-odds for a unit change in x

Classification with Logistic Regression: Bayesian

```
with pm.Model() as model_lrs:
    # Linear part.
    alpha = pm.Normal("alpha", mu=0, sigma=1)
    beta = pm.Normal("beta", mu=0, sigma=5)
    mu = alpha + x_c * beta
    # Sigmoid.
    theta = pm.Deterministic("theta", pm.math.sigmoid(mu))
    # Model.
    yl = pm.Bernoulli("yl", p=theta, observed=y_0)
    # Intercept?
    bd = pm.Deterministic("bd", - alpha / beta)
    #
    idata_lrs = pm.sample(random_seed=123)
```


Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [alpha, beta]

Sampling 4 chains, 0 divergences



Heteroskedasticity



- **Heteroskedasticity** is when variance of errors is not constant across observations
 - E.g., variance as a linear function of the dependent variable
- **Example:**
 - “Baby height as a function of age” is heteroskedastic
 - Mean μ as square root of a linear model
 - Standard deviation σ as a linear function of the predictor variable



Heteroskedasticity: Bayesian Model

```
with pm.Model() as model_vv:
    # Create a shared variable so that the data can change after the model is created.
    x_shared = pm.Data("x_shared", data.month.values.astype(float))
    # Linear model for the mean is a function of sqrt(x).
    alpha = pm.Normal("alpha", sigma=10)
    beta = pm.Normal("beta", sigma=10)
    mu = pm.Deterministic("mu", alpha + beta * x_shared ** 0.5)
    # Linear model for the std dev.
    gamma = pm.HalfNormal("gamma", sigma=10)
    delta = pm.HalfNormal("delta", sigma=10)
    sigma = pm.Deterministic("sigma", gamma + delta * x_shared)
    # Fit.
    y_pred = pm.Normal("y_pred", mu=mu, sigma=sigma, observed=data.length)
    #
    idata_vv = pm.sample(random_seed=123)
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [alpha, beta, gamma, delta]
```



- Generalized Linear Models
 - Simple Linear Model
 - Logistic Regression
 - *Multiple Linear Regression*

Multiple Linear Regression

- In prediction problems, common to use **several independent variables**
 - E.g., student's grades = $f(\text{family income, mother's education, } \dots)$
- **Problem formulation**
 - k independent variables
 - N observations
 - Find a hyperplane of dimension k to explain data

$$\mu = \alpha + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k$$

- Similar to polynomial regressions with independent variables

Multiple Regression: Synthetic Example 1/2


■ Generate random data

```
np.random.seed(314)

N = 100
#N = 1000
alpha_real = 2.5
beta_real = [0.9, 1.5]
eps_stddev_real = 0.5
eps_real = np.random.normal(0, eps_stddev_real, size=N)

# Independent variables.
X = np.array([np.random.normal(i, j, N) for i, j in zip(
    # mean of gaussian.
    [10, 2],
    # std dev.
    [1, 1.5])]).T
X_mean = X.mean(axis=0, keepdims=True)
X_centered = X - X_mean

# Create samples.
y = alpha_real + np.dot(X, beta_real) + eps_real
```




■ PyMC model

```
with pm.Model() as model_mlr:
    alpha_tmp = pm.Normal('alpha_tmp', mu=0, sigma=10)
    # Beta is a vector.
    beta = pm.Normal('beta', mu=0, sigma=1, shape=2)
    eps = pm.HalfCauchy('eps', 5)
    # mu.
    mu = alpha_tmp + pm.math.dot(X_centered, beta)
    # Extract alpha.
    alpha = pm.Deterministic('alpha', alpha_tmp - pm.math.dot(X_mean, beta))

    # Model.
    y_pred = pm.Normal('y_pred', mu=mu, sigma=eps, observed=y)


    idata_mlr = pm.sample(2000)
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [alpha_tmp, beta, eps]
```



Multiple Regression: Synthetic Example 2/2

- Solve model



	mean	sd	hdi_3%	hdi_97%
alpha[0]	1.85	0.46	1.01	2.75
beta[0]	0.97	0.04	0.88	1.05
beta[1]	1.47	0.03	1.41	1.53
eps	0.47	0.03	0.41	0.54


Multiple Regression: Rented Bike Example 1/2

- **Assumption:** number of bike rented is function of temperature and hour of the day
- PyMC model

```
with pm.Model() as model_mlb:
    alpha = pm.Normal("alpha", mu=0, sigma=1)
    beta0 = pm.Normal("beta0", mu=0, sigma=10)
    beta1 = pm.Normal("beta1", mu=0, sigma=10)
    sigma = pm.HalfNormal("sigma", 10)
    mu = pm.Deterministic("mu", pm.math.exp(alpha + beta0 * bikes.temperature +
    beta1 * bikes.hour))
    _ = pm.NegativeBinomial("y_pred", mu=mu, alpha=sigma, observed=bikes.rented)
    #
    idata_mlb = pm.sample()
```


```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [alpha, beta0, beta1, sigma]
```

```
Sampling 4 chains, 0 divergences
```



Multiple Regression: Rented Bike Example 2/2

- Solve model



	mean	sd	hdi_3%	hdi_97%
alpha	3.55	0.13	3.32	3.82
beta0	0.04	0.01	0.03	0.05
beta1	0.08	0.01	0.06	0.10
sigma	1.11	0.08	0.97	1.25

Tutorial

- [Generalized Linear Models](#)